

# APR - Aprendizagem Profunda e por Reforço - 2025/26

Deep and Reinforcement Learning

## P31 – CNN & CNN

**Jorge Henriques, João Nuno Correia**

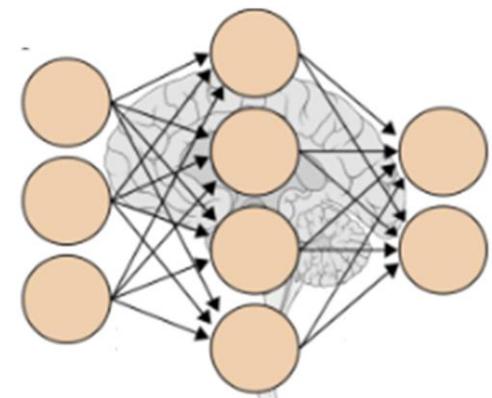
[jh@dei.uc.pt](mailto:jh@dei.uc.pt), [jncor@dei.uc.pt](mailto:jncor@dei.uc.pt)

Departamento de Engenharia Informática  
Faculdade de Ciências e Tecnologia



UNIVERSIDADE  
DE  
COIMBRA

**dei** engenharia  
informática



## ■ Main Goals

- Use our library on other problems (e.g., the Iris dataset)
- Introduction to the PyTorch library
  
- PyTorch to define and train a Multilayer Neural Network  
**(MLN)**
- PyTorch to define and train a Convolutional Neural Network  
**(CNN)**

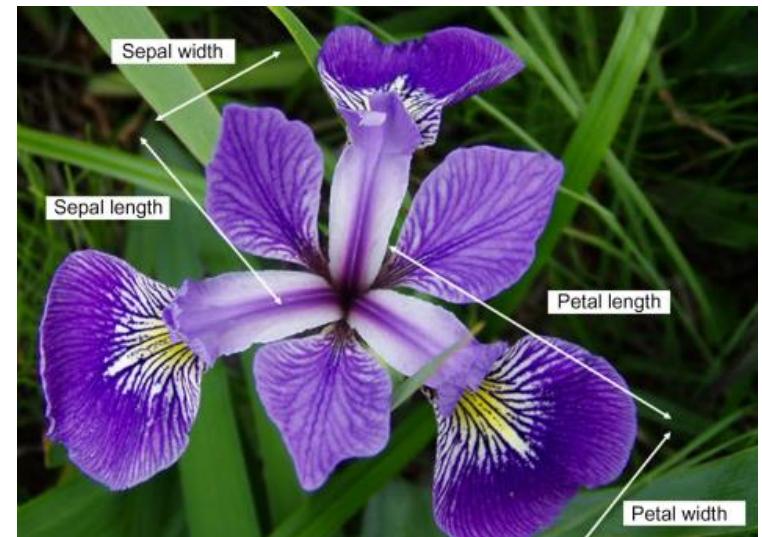
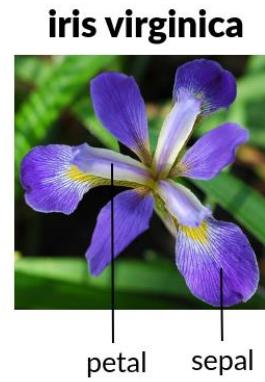
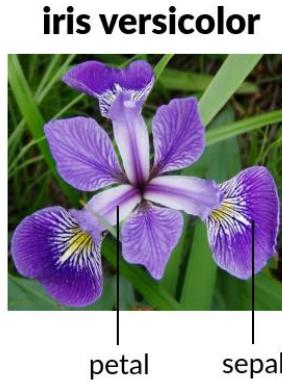
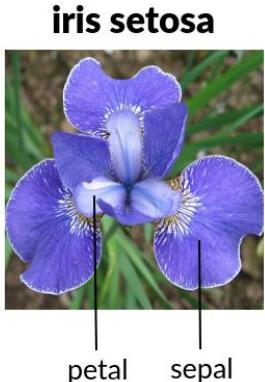
# Contents

- 1 | Iris Dataset
- 2 | Pytorch
- 3 | MLNN to MINIST dataset
- 4 | MLNN to CIFAR-10 dataset
- 5 | CNN to CIFAR-10 dataset

## Iris dataset

```
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
Y = iris.target
```

- N=150
- Features: 4 (sepal length, sepal width, petal length, petal width)
- Classes: 3 (Setosa, Versicolor, Virginica)



## Machile Learning / Deep learning models

### Manual feature extraction

sepal length,  
sepal width,  
petal length,  
petal width

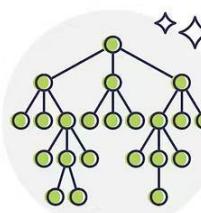
### MACHINE LEARNING



INPUT

FEATURE EXTRACTION

CLASSIFICATION



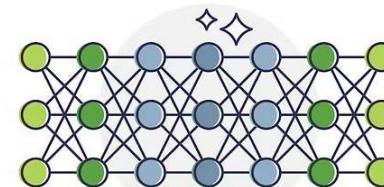
**iris virginica**

### DEEP LEARNING



INPUT

FEATURE EXTRACTION + CLASSIFICATION



**iris virginica**

- Dataset: train and test

```
X_train, X_test, y_train, y_test = train_test_split(X, Y,  
test_size=0.2, random_state=10)
```

- Splits the dataset into two parts: **Training set** (80%), **Test set** (20%)

- $0.8 \times 150 = 120$

*X\_train (120,4)*

*Y\_train (120,1)*

- $0.2 \times 150 = 30$

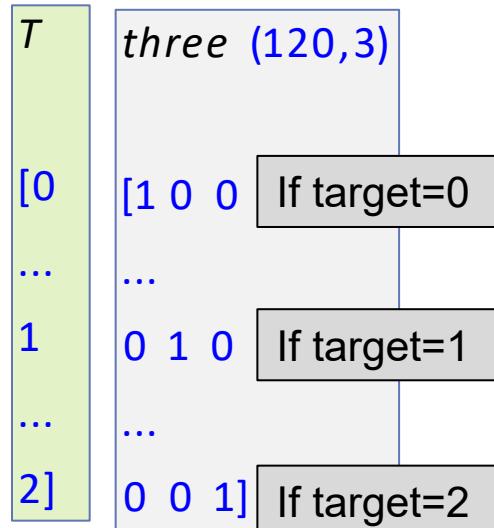
*X\_test (30,4)*

*Y\_test (30,1)*

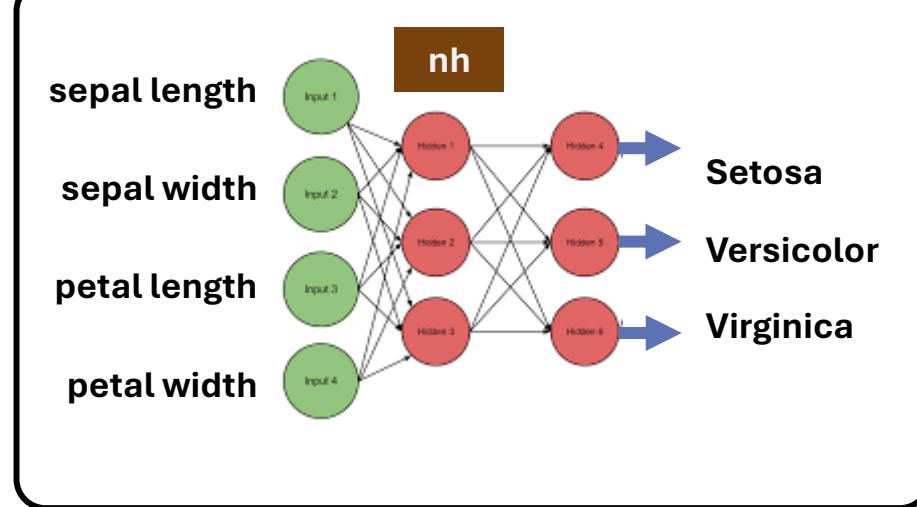
## ■ Classes

```
three = np.zeros((y_train.shape[0],3))      # (120,3)
three[np.where(y_train == 0),0] = 1
three[np.where(y_train == 1),1] = 1
three[np.where(y_train == 2),2] = 1
```

- Output 1D to classes 3D



Neural network structure



## ■ Normalization

```
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
scaler.fit(X_train)  
scaler.transform(X_train);  
X_train = np.array(scaler.transform(X_train))  
X_test = np.array(scaler.transform(X_test))
```

- StandardScaler is used for feature scaling, which standardizes features by removing the mean and scaling to unit variance.
- Returns a scaled array where each **feature has mean 0 and variance 1**.

$$X_{scaled} = \frac{X - \mu}{\sigma}$$

## Define and train a MLNN

Inputs=4, outputs=3

Nh=30

```
an = MultiLayerPerceptron(4, 3, 30, number_of_iterations=1000,
    output_type='logistic', learning_rate = 0.001)

an.train(X_train, three)
```

```
Y= an.predict (X_test)
test_predictions = np.argmax(Y, axis=1)      # {0,1,2}

print("Confusion Matrix\n", confusion_matrix(y_test, test_predictions))
print("Accuracy: %.2f %%" % (accuracy_score(y_test, test_predictions) * 100))
```

Confusion Matrix  
[[10 0 0]  
 [ 0 13 0]  
 [ 0 0 7]]  
Accuracy: 100.00 %

# Contents

- 1 | Iris Dataset
- 2 | Pytorch
- 3 | MLNN to MINIST dataset
- 4 | MLNN to CIFAR-10 dataset
- 5 | CNN to CIFAR-10 dataset

 PyTorch

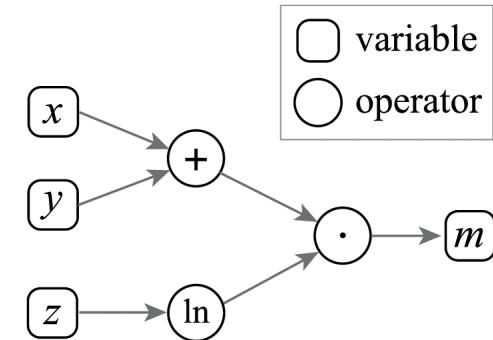
- PyTorch is an open-source machine learning library primarily developed by Facebook's AI Research lab (FAIR).
- Popular among researchers and developers for building and training deep learning models.
- This feature is particularly useful for tasks such as **natural language processing and computer vision**.



## Key features of PyTorch include:

- **Tensors:**
  - Multi-dimensional array called a tensor,
  - Similar to NumPy arrays
  - Added benefit of **GPU acceleration** for numerical computations.
- **Neural Network Module:**
  - `torch.nn` module, providing pre-defined layers, loss functions, and optimization algorithms to facilitate the construction of neural network architectures
- **Autograd:**
  - Automatic differentiation library called Autograd, which **automatically computes gradients for tensors during backpropagation.**
  - This simplifies the process of training neural networks.

# PyTorch



- A **computational graph** is a structure that represents **all the operations and dependencies** in a neural network.
- PyTorch builds its **computational graph dynamically**,
  - It creates the graph as your code runs.
  - This is called “define-by-run”: you don’t have to define the whole network in advance.
  - **Easy debugging:** You can use normal Python tools like `print()` to inspect what’s happening.
  - **Flexible models:** You can write loops, conditionals, or work with inputs of different sizes.

## ■ 1. Creating Tensors

```
import torch

x = torch.tensor([1.0, 2.0, 3.0])
x = torch.zeros(2, 3)

x = torch.tensor([1.0, 2.0, 3.0], requires_grad=True)*

# Convert numpy to a PyTorch tensor
np_array = np.array([[1, 2], [3, 4]])
tensor = torch.from_numpy(np_array)

* requires_grad=True → PyTorch will compute gradients for this tensor.
```

## ■ 2. Basic Tensor Operations

```
a = torch.tensor([1.0, 2.0])
b = torch.tensor([3.0, 4.0])
c = a + b
d = a * b

#----- Mean, sum
mean_val = a.mean()
sum_val = b.sum()

#----- # Matrix multiplication
x = torch.rand(2, 3)
y = torch.rand(3, 2)
z = torch.matmul(x, y)
```

## ■ 3. Using GPU (CUDA)

```
if torch.cuda.is_available():
    device = torch.device("cuda:0")
else:
    device = torch.device("cpu")

#----- Move tensors to device
x = x.to(device)

#----- Move a model to device
model.to(device)
```

CUDA Compute Unified Device Architecture.

Parallel computing platform and API created by NVIDIA that lets software (PyTorch, TensorFlow, etc.) use the GPU (Graphics Processing Unit) to do computations

## ■ 4. Neural networks: definitions

```
import torch
import torch.nn as nn # Neural network modules
import torch.optim as optim # Optimizers

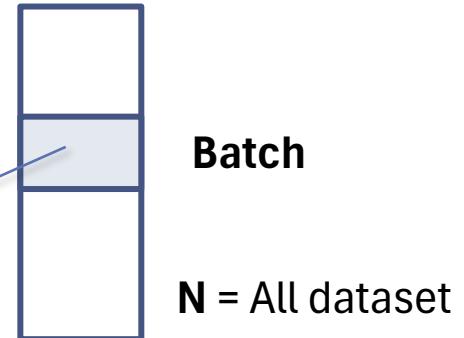
#----- Loss functions

criterion = nn.MSELoss() # For regression
criterion = nn.CrossEntropyLoss() # For classification

#----- Optimizers

optimizer = optim.SGD(model.parameters(), lr=0.01)
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

## ■ 5. Neural networks: Train Loop



```

Num_epochs=100

for epoch in range(num_epochs):
    for X_batch, y_batch in train_loader:
        # iterates over all epochs
        # Over the dataset in mini-batches
        # using a Data Loader

        X_batch = X_batch.to(device)           # To device=CPU/GPU
        y_batch = y_batch.to(device)

        outputs = model(X_batch)             # Step 1: Forward pass
        loss = criterion(outputs, y_batch)   # Step 2: Compute loss

        optimizer.zero_grad()                # Step 3: Clear old gradients
        loss.backward()                     # Step 4: Backpropagation
        optimizer.step()                   # Step 5: Update weights
    
```

# Contents

- 1 | Iris Dataset
- 2 | Pytorch
- 3 | **MLNN to MINIST dataset**
- 4 | MLNN to CIFAR-10 dataset
- 5 | CNN to CIFAR-10 dataset

0	8	7	6	4	6	9	7	2	1	5	1	4	6	
0	1	2	3	4	4	6	2	9	3	0	1	2	3	4
0	1	2	3	4	5	6	7	0	1	2	3	4	5	0
7	4	2	0	9	1	2	8	9	1	4	0	9	5	0
0	2	7	8	4	8	0	7	7	1	1	2	9	3	6
5	3	9	4	2	7	2	3	8	1	2	9	8	8	7
2	9	1	6	0	1	7	1	1	0	3	4	2	6	4
7	7	6	3	6	7	4	2	7	4	9	1	0	6	8
2	4	1	8	3	5	5	5	3	5	9	7	4	8	5

## ■ Load Data → numpy

0	8	7	6	4	6	9	7	2	1	5	1	4	6
0	1	2	3	4	4	6	2	9	3	0	1	2	3
0	1	2	3	4	5	6	7	0	1	2	3	4	5
7	4	2	0	9	1	2	8	9	1	4	0	9	5
0	2	7	8	4	8	0	7	7	1	1	2	9	3
5	3	9	4	2	7	2	3	8	1	2	4	8	8
2	9	1	6	0	1	7	1	1	0	3	4	2	6
7	7	6	3	6	7	4	2	7	4	9	1	0	6
2	4	1	8	3	5	5	5	3	5	7	4	8	5

```
#----- LOAD Train/test data
#----- as a tensor
# Train = 60000, 28x28, Test = 10000 28x28

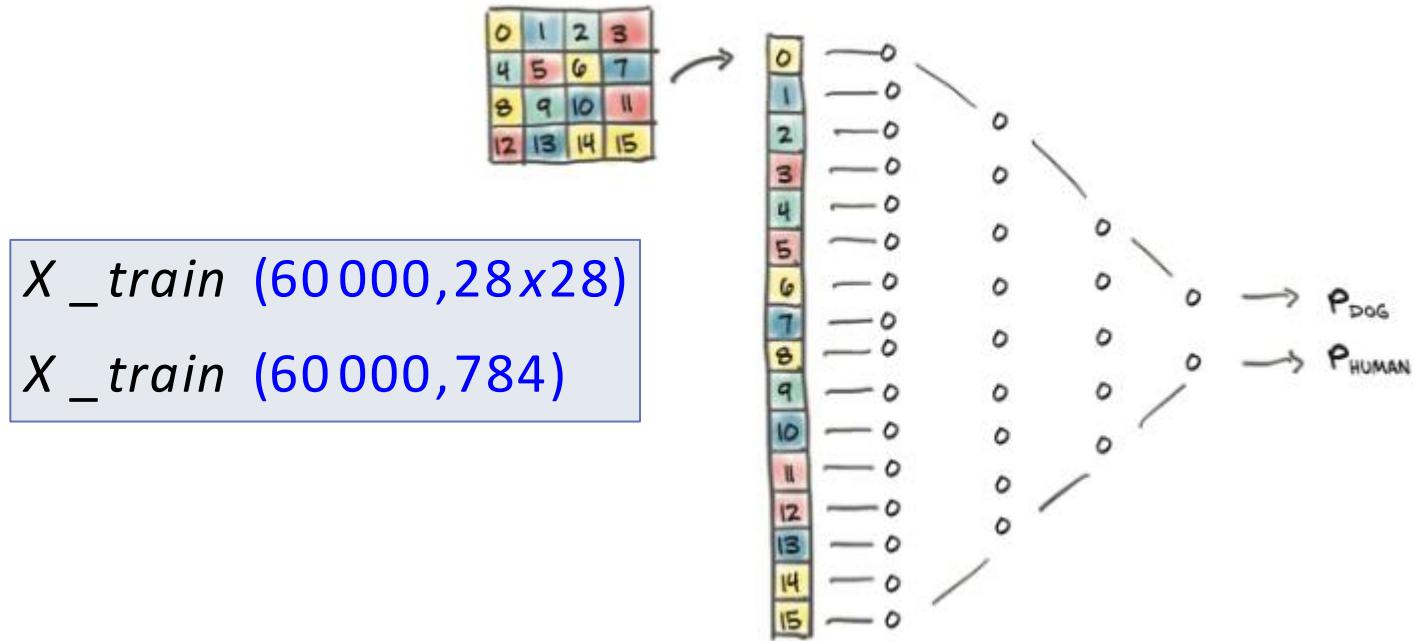
mnist_trainset =
datasets.MNIST(root='./data', train=True, download=True, transform=None)
mnist_testset =
datasets.MNIST(root='./data', train=False, download=True, transform=None)

#----- Tensor to numpy
X_train = mnist_trainset.data.numpy()
y_train = mnist_trainset.targets.numpy()
X_test = mnist_testset.data.numpy()
y_test = mnist_testset.targets.numpy()
```

## ■ Binary images (black and white) & Flattening

```
#----- Image [0 .. 255]
#----- Convert to [0..1]
X_train = np.where(X_train>0.5,1,0)
X_test = np.where(X_test>0.5,1,0)

# Flatten the input (i.e., convert 2D images to 1D vectors)
X_train = X_train.reshape(X_train.shape[0], -1)
X_test = X_test.reshape( X_test.shape[0], -1)
```



- Data to tensor &
- Number of inputs and outputs=classes

```
#----- Image [0 .. 255]
X_train = torch.tensor(X_train).float()
y_train = torch.tensor(y_train).long()
X_test = torch.tensor( X_test).float()
y_test = torch.tensor( y_test).long()

num_inputs = X_train.shape[1]                      # 784
num_classes = int(y_train.max().item() + 1)         # 10 classes
```

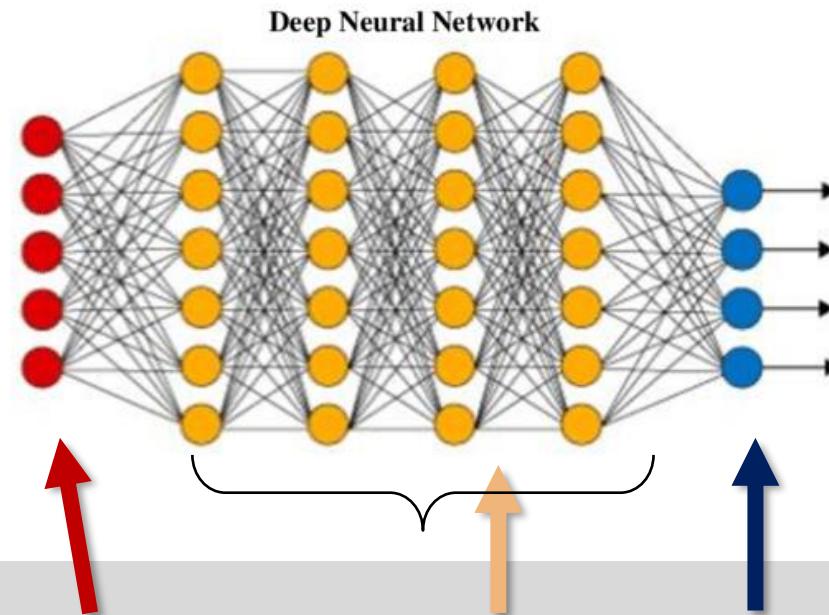
## ■ Device

```
if torch.backends.mps.is_available():      # **
    device = torch.device("mps")

elif torch.cuda.is_available():
    device = torch.device("cuda")
else:
    device = torch.device("cpu")
```

**# MPS -- Apple's Metal Performance Shaders**

## ■ MLNN - Deep CNN



```
class DNN(nn.Module):
    def __init__(self, input_size, hidden_sizes, num_classes):
        self.layers = nn.ModuleList()
        self.activations = nn.ModuleList()

# ----- Example
NumInp= 5
NHidden=[7,7,7] # same size !!!
NumCla= 4

dnn = DNN(input_size=NumInp, hidden_sizes=NHidden, num_classes=NumCla)
```

## ■ TODO

```
self.layers = nn.ModuleList()          # from torch
self.activations = nn.ModuleList()

self.layers      is a list of linear layers, e.g. nn.Linear
self.activations is a list of activation functions, e.g. nn.ReLU()
```

```
def forward(self, x):
    ----- TODO: implement the forward pass
    out = x
    for i in range(len(self.layers)):
        out = self.layers[i](out)           # All layers
        if i < len(self.activations):
            out = self.activations[i](out)  # out = activation(out)

    return out
```

## ■ Train the MLNN

```
def fit(X_train, y_train, nn, criterion, optimizer, n_epochs,
       to_device=True, batch_size=32):
    ...
    return loss_values, nn

# X_train           Input training data
# y_train           Target training data
# nn               DNN – MLNN
# criterion        Loss functions
# optimizer        Backpropagation optimizer
# n_epochs         Number of epochs
# batch_size       Dimension of each batch for training
```

## ■ Application to MNIST

```

num_inputs = 28*28
num_classes = 10
n_epochs = 100
n_layers = 3
BATCH_SIZE = 100

# Determine the number of hidden_layer_sizes = [397,397,397,397]
hidden_layer_sizes = ((num_inputs + num_classes) // 2,) * n_layers
# MUST HAVE THE SAME SIZE !!

#----- Define the MLNN
dnn =
DNN(input_size=num_inputs, hidden_sizes=hidden_layer_sizes, num_classes=num_classes)

#----- Definitions
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(dnn.parameters(), lr=0.1)

#----- Train MLNN
loss_values, dnn =
fit(X_train, y_train, dnn, criterion, optimizer, n_epochs, batch_size=BATCH_SIZE,
to_device=False)

```

$$nh = \frac{Numinputs + NumClasses}{2}$$

$$nh = \frac{784 + 10}{2} = 397$$

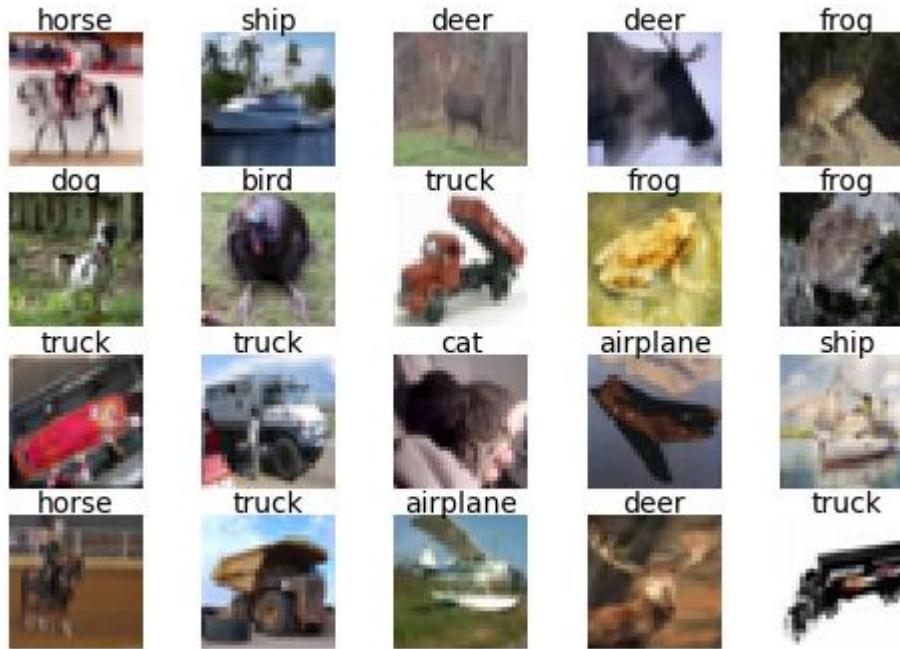
# MUST HAVE THE SAME SIZE !!

# Contents

- 1 | Iris Dataset
- 2 | Pytorch
- 3 | MLNN to MINIST dataset
- **4 | MLNN to CIFAR-10 dataset**
- 5 | CNN to CIFAR-10 dataset

## CIFAR-10

- Similar to MNIST,
- Classification problem with 10 possible classes.
- Images are more complex, representing real life scenes instead of hand-written digits.



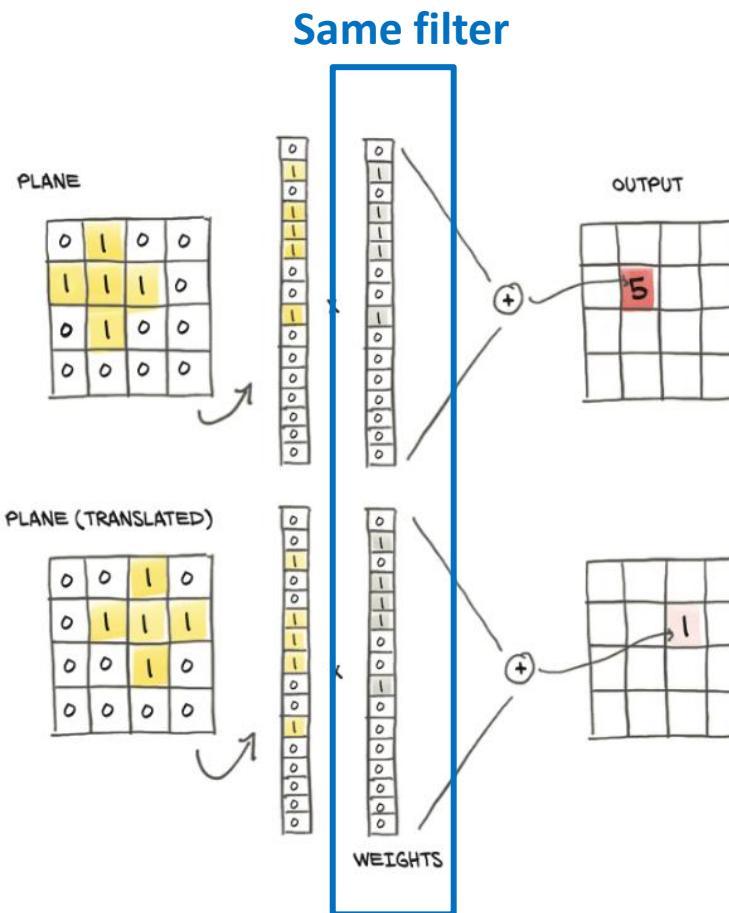
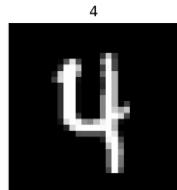
0 - airplane  
1 - automobile  
2 - bird  
3 - cat  
4 - deer  
5 - dog  
6 - frog  
7 - horse  
8 - ship  
9 - truck

## MLNN

- Main problem !!
- The use of pixel value values make the model **sensible to translations**



Same pattern

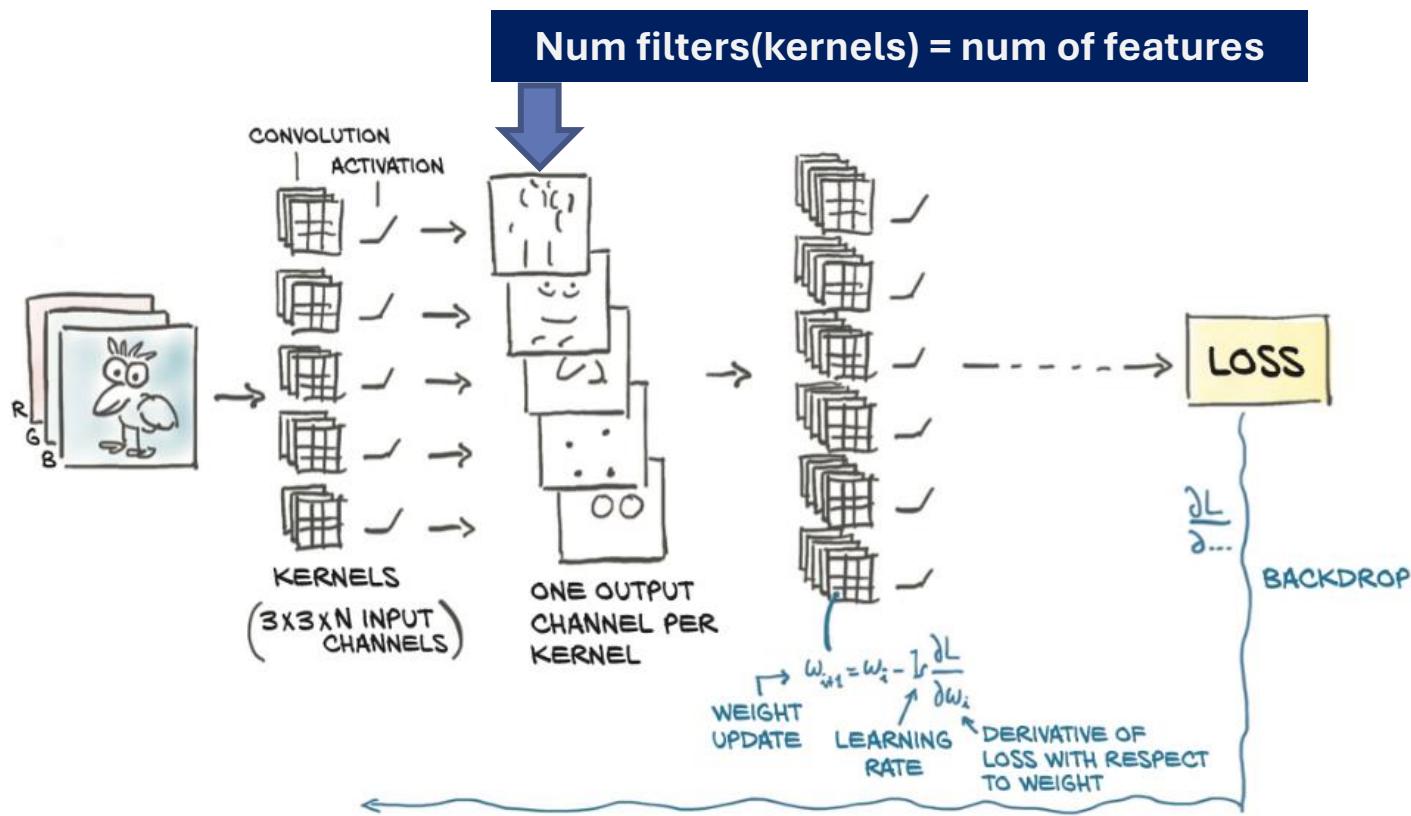


# Contents

- 1 | Iris Dataset
- 2 | Pytorch
- 3 | MLNN to MINIST dataset
- 4 | MLNN to CIFAR-10 dataset
- 5 | CNN to CIFAR-10 dataset

## CNN – Convolutional neural networks

- CNN uses these operations to **automatically compute features** on an input image producing feature maps.
- Translation-invariant characteristics, no matter the variations in position, orientation, scale, or translation
- Filters (weights) are adjusted to minimize a loss function**



## Build a CNN

```

class CNN(nn.Module):
    def __init__(self, input_channels=3, num_classes=10):
        #----- Two Convolutional layers
        self.conv1 = nn.Conv2d(input_channels, 16, kernel_size=3, padding=1) 32x32
        self.conv2 = nn.Conv2d(16, 8, kernel_size=3, padding = 1) 16 → pool(2x2)=8
        #----- Two fully connect layers
        self.fc1 = nn.Linear(8*8*8, 32)
        # 8x8x8=512 (height × width × channels) the tensor after the last convolution
        self.fc2 = nn.Linear(32, num_classes)

```

**RBG**

**16 filters**

**3x3**

**8 filters**

**3x3**

myCNN=CNN(3,10)

## ■ Train a CNN

```
cnn_loss_values, cnn =  
    fit(X_train, y_train, cnn, criterion, optimizer, n_epochs, to_device=True,  
        batch_size=BATCH_SIZE)
```